

SQL

What is SQL?



Structured Query Language

A querying language designed for accessing and manipulating information from Relational Databases

SQL


Advantages



Platform
Independent




Portable



Database
Independent



Simple
Syntax



Quick and Efficient
Retrieval

SQL

Advantages

- SQL is portable and can be used independent of the platform
- Can be used for querying data in a wide variety of databases and data repositories,
- Has a simple syntax that is similar to the english language
- Can retrieve large amounts of data quickly and efficiently
- Runs on an interpreter system (which means code can be executed as soon as it is written, making prototyping quick and easy)

SQL

Command Types



DQL
(Data
Query
Language)

SELECT



DDL
(Data
Definition
Language)

**CREATE
DROP
ALTER
RENAME**



DML
(Data
Manipulation
Language)

**INSERT
UPDATE
DELETE**



TCL
(Transaction
Control
Language)

**COMMIT
ROLLBACK
SAVEPOINT**



DCL
(Data
Control
Language)

**GRANT
REVOKE**

SQL

Command Types

- DDL: Data Definition Language
 - Define admissible database content (schema)
- DQL: Data Query Language
 - Query and retrieve database content
- DML: Data Manipulation Language
 - Change and retrieve database content
- TCL: Transaction Control Language
 - Groups SQL commands (transactions)
- DCL: Data Control Language
 - Assign data access rights

DQL

SQL

DQL



DQL
Data Query Language




Query and retrieve data from database

Describes a new relation to generate



DQL
(Data
Query
Language)

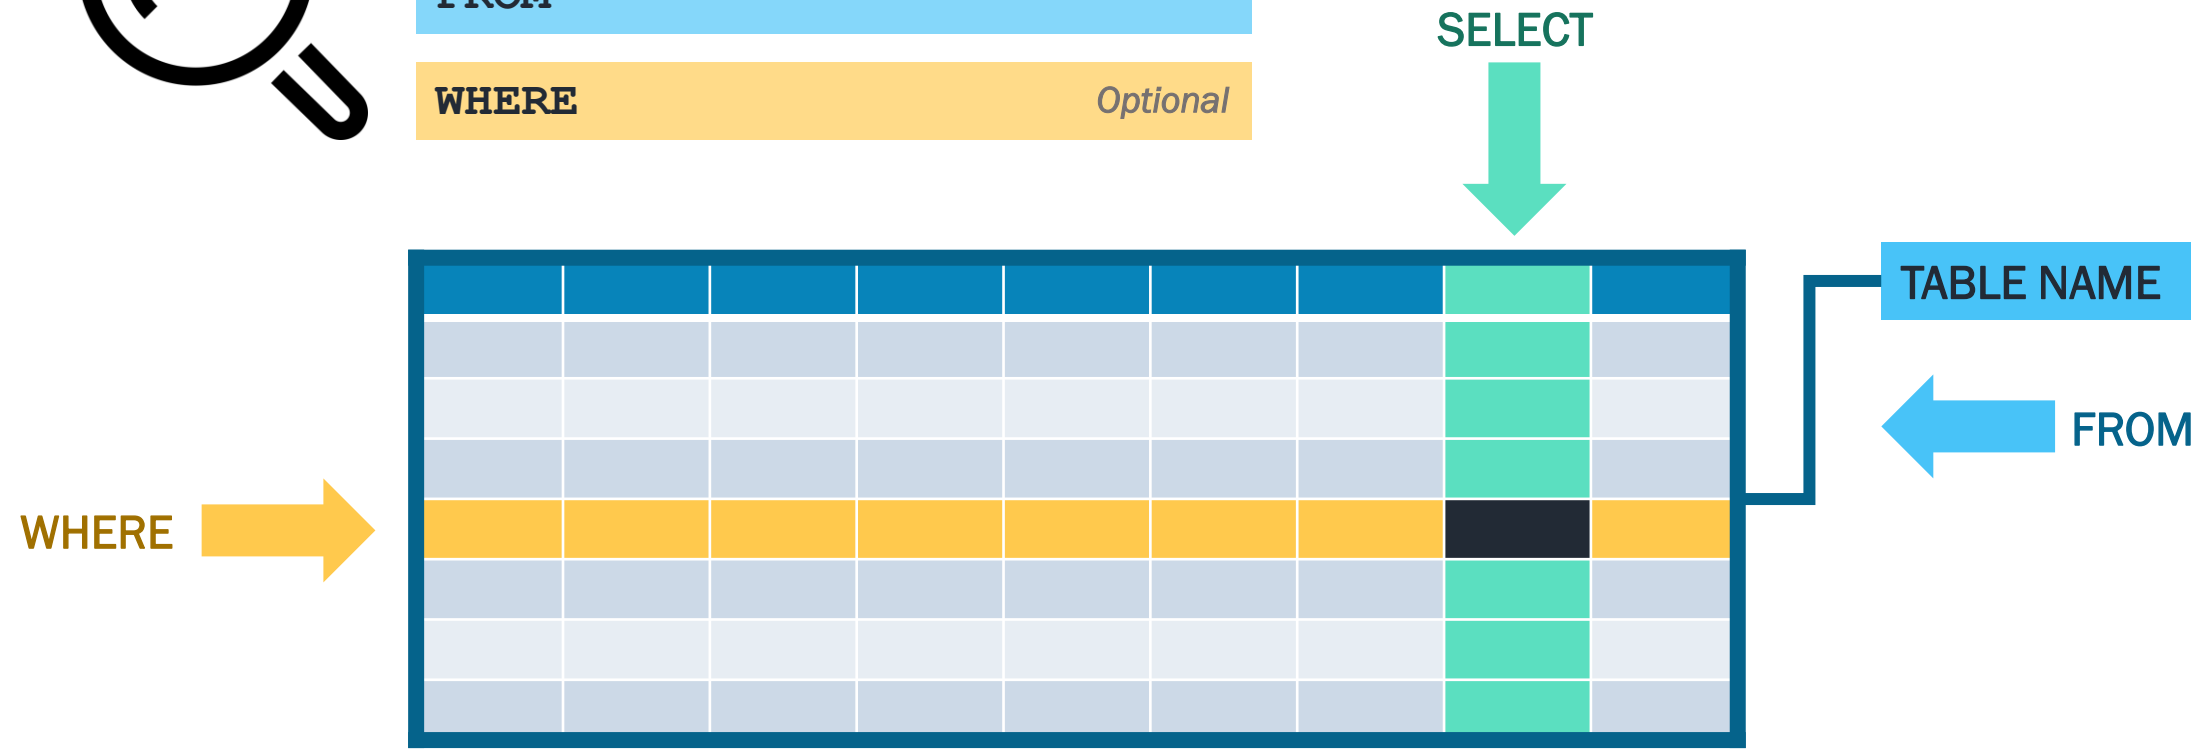
SELECT



A simple SQL Query consists of 3 main clauses:



- SELECT
- FROM
- WHERE *Optional*



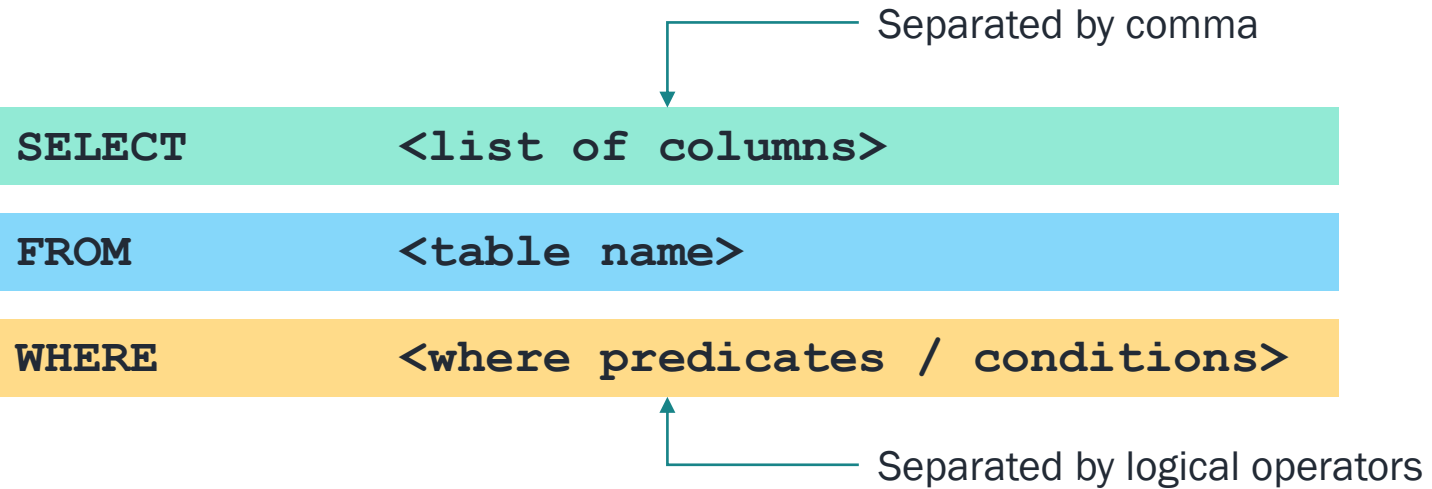
SQL

DQL

- What is an SQL query? A SQL query describes a new relation to generate
- A simple SQL query consist of 3 clauses:
 - SELECT: describes the columns of relation to generate
 - FROM: describes source relations and how to match
 - WHERE: defines conditions result rows must satisfy
- The syntax is:
 - SELECT <column_list>
 - FROM <table1_name> JOIN <table2_name> ON (<join_predicate>)
 - WHERE <where_predicate>
- where <column_list> is a comma separated list of columns and <table1_name> and <table2_name> are database relations, <join_predicate> is a condition defining matching tuples pairs and <where_predicate> are additional conditions.

SQL

DQL



Example:

```
SELECT      first_name, last_name
FROM        student_table
WHERE       student_id = 13787545
```

SQL

DQL : *



```
SELECT *
```

Selects (retrieves) all the columns from a table

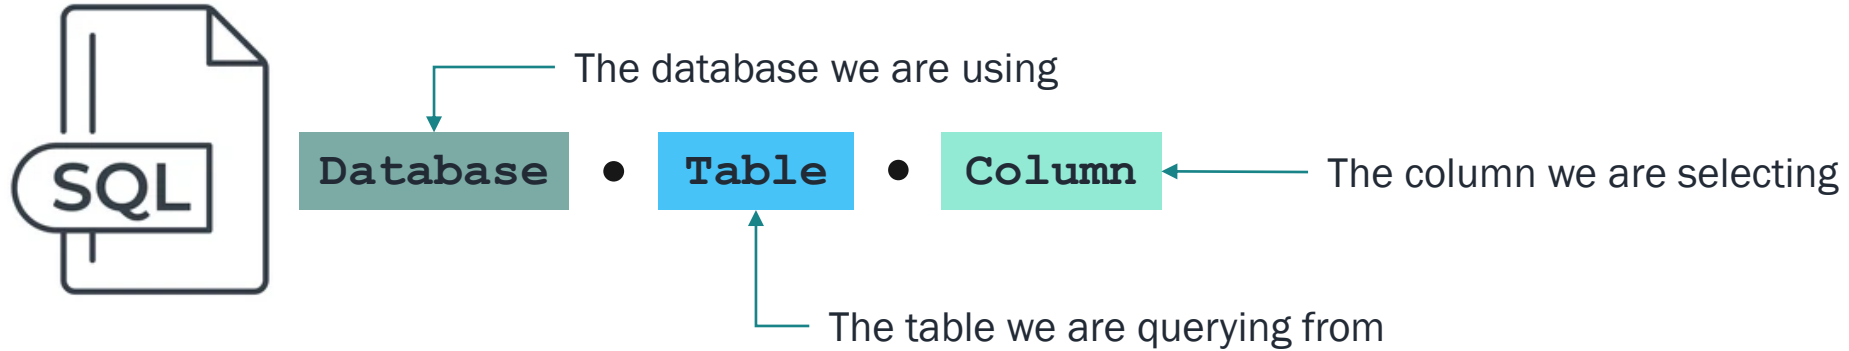
Example:

```
SELECT *  
FROM customers
```

SQL

DQL : .

A dot (.) can be used to specify objects in hierarchical form



Example:

```
SELECT      coffeeshop.customers.customer_id
FROM        coffeeshop.customers
```

SQL

Aliases

Aliases are used to give a column a temporary name while in a query



```
SELECT      column1 AS c1
```

Example:

```
SELECT      loyalty_card_number AS lcn  
FROM        customers
```

SQL

Aliases

Aliases are also used to give a table a temporary name while in a query



```
FROM      table1 AS t1
```

Example:

```
SELECT      *  
FROM      customers AS c
```

```
SELECT      c.customer_id  
FROM      customers AS c
```

SQL

DQL: Distinct

Used to return only distinct (different) values (to eliminate duplicates)



```
SELECT      DISTINCT Column1
```

Example:

```
SELECT      DISTINCT continent  
FROM        country
```

```
SELECT      DISTINCT GovernmentForm  
FROM        country
```


SQL

DQL: Order By

Used to sort the outcome table by a set of columns



```
SELECT      <list of columns>
```

```
FROM        <table name>
```

```
ORDER BY    <list of columns>
```

By default, is in Ascending order
Use **DESC** to sort Descendingly

Example:

```
SELECT      *  
FROM        country  
ORDER BY    population DESC
```

SQL

DQL: Limit

Used to limit the number of rows returned



```
SELECT      <list of columns>
```

```
FROM        <table name>
```

```
LIMIT      <number>
```

Example:

```
SELECT      *  
FROM        country  
LIMIT      10
```

Operators

SQL

Operators

Arithmetic Operators



Addition, Subtraction, Multiplication, Division, Modulus

Example:

```
SELECT      current_wholesale_price,  
            3*(current_wholesale_price+2) AS new_price  
FROM        coffeeshop.products;
```

```
SELECT      *  
FROM        products  
WHERE       current_wholesale_price + 3 < 10
```

SQL

Operators

Comparison Operators



= > < >= <= <> !

Used to compare the values of two operands

Example:

```
SELECT *  
FROM products  
WHERE tax_exempt_yn = 'Y'
```

```
SELECT *  
FROM products  
WHERE current_wholesale_price + 3 < 10
```

SQL

Operators

Logical Operators



AND

TRUE if all the conditions separated by AND is TRUE

Condition 1	AND	Condition 2	>	Result
TRUE		TRUE		TRUE
TRUE		FALSE		FALSE
FALSE		FALSE		FALSE

SQL

Operators

Logical Operators



OR

TRUE if any of the conditions separated by OR is TRUE

Condition 1	OR	Condition 2	>	Result
TRUE		TRUE		TRUE
TRUE		FALSE		TRUE
FALSE		FALSE		FALSE

SQL

Operators

Logical Operators



NOT

Used to negate (reverse) the meaning of a comparison or a logical operator

Examples:

NOT	a > b	≡	a <= b	or	! a > b
NOT	a = b	≡	a != b	or	a <> b

NOT	Condition	>	Result
NOT	TRUE		FALSE
NOT	FALSE		TRUE

SQL

Unknown Values

NULL Values



NULL

Represents missing, unknown, or blank

Arithmetic with NULL values:

NULL + - * / % anything = NULL

Comparison with NULL values:

NULL = > < >= <= <> anything ≡ NULL

SQL

Operators

Logical Operators



IS

IS NOT

Used to check if a value is NULL or not

Examples:

anything

IS

NULL

≡

TRUE

or

FALSE

anything

IS NOT

NULL

≡

TRUE

or

FALSE

Exercise: Can these expressions ever result in NULL? Why?

SQL

Three-Valued Logic

Ternary Outcomes



TRUE

FALSE

NULL

Outcome of an expression can be evaluated as either TRUE, FALSE, or NULL

Condition 1

OR

Condition 2

TRUE

NULL

FALSE

NULL

Condition 1

AND

Condition 2

TRUE

NULL

FALSE

NULL

Exercise: What are outcomes of these expressions?

SQL

Operators - Exercise

What are the outcomes of the following expressions?

1	SELECT	4 = NULL
2	SELECT	NULL = NULL
3	SELECT	NULL IS NULL
4	SELECT	NULL IS NOT NULL
5	SELECT	TRUE OR NULL
6	SELECT	FALSE OR NULL
7	SELECT	TRUE AND NULL
8	SELECT	FALSE AND NULL
9	SELECT	(TRUE AND FALSE) OR NULL

SQL

Operators - Answer

What are the outcomes of the following expressions?

1	SELECT	4 = NULL	NULL
2	SELECT	NULL = NULL	NULL
3	SELECT	NULL IS NULL	TRUE
4	SELECT	NULL IS NOT NULL	FALSE
5	SELECT	TRUE OR NULL	TRUE
6	SELECT	FALSE OR NULL	NULL
7	SELECT	TRUE AND NULL	NULL
8	SELECT	FALSE AND NULL	FALSE
9	SELECT	(TRUE AND FALSE) OR NULL	NULL

SQL

Operators

Logical Operators



AND

OR

Example:

```
SELECT *  
FROM country  
WHERE LifeExpectency > 75 OR Population > 50000000
```

```
SELECT *  
FROM country  
WHERE IndepYear > 1950 AND GNP > 1000000
```

```
SELECT *  
FROM country  
WHERE Population < 100000000 AND GNP > 1000000
```

SQL

Operators

Logical Operators



BETWEEN

TRUE if the operand is within the range of comparisons

Example:

```
SELECT      *  
FROM        country  
WHERE       LifeExpectency BETWEEN 50 AND 60
```

SQL

Operators

Logical Operators



IN

TRUE if the operand is equal to one of a list of expressions

Example:

```
SELECT *  
FROM country  
WHERE Continent IN ('Africa', 'Asia')
```

```
SELECT *  
FROM country  
WHERE IndepYear IN (1990, 1991, 1992)
```


SQL

Operators

Logical Operators & Regular Expressions



LIKE

TRUE if the operand matches a pattern

Wild Cards used:

_

%

Matches any number of characters

Represents a single character

Example:

```
SELECT *  
FROM country  
WHERE Continent LIKE 'A%'
```

```
SELECT *  
FROM country  
WHERE Name LIKE 'C%'
```

Logical Operators & Regular Expressions



LIKE

TRUE if the operand matches a pattern

Examples:

column LIKE 'a%'	Finds any values that start with "a"
column LIKE '%a'	Finds any values that end with "a"
column LIKE '%or%'	Finds any values that have "or" in any position
column LIKE '_r%'	Finds any values that have "r" in the second position
column LIKE 'a_%'	Finds any values that start with "a" and are at least 2 characters in length
column LIKE 'a__%'	Finds any values that start with "a" and are at least 3 characters in length
column LIKE 'a%o'	Finds any values that start with "a" and ends with "o"

Aggregation

SQL

Aggregation

Simple Aggregation



SUM

MIN

MAX

AVG

STD

COUNT

Aggregates all the records of a column by taking the SUM, MIN, MAX, AVG, STD or the number of records.

Example:

```
SELECT      AVG (LifeExpectancy)
FROM        country;
```

```
SELECT      COUNT (*)
FROM        country
WHERE       Continent = 'Asia'
```

SQL

Aggregation

Group By Aggregation Used to summarize data based on values of a specific column



```
SELECT      Agg Func (<column name>)
```

```
FROM        <table name>
```

```
WHERE       <where predicates / conditions>
```

```
GROUP BY    <list of columns>
```

Example:

```
SELECT      continent, AVG(LifeExpectancy)
FROM        country
WHERE       Population > 30000
GROUP BY    Continent
```

SQL

Aggregation

Group By Aggregation with Conditions Used to apply a condition to the groups



SELECT Agg Func (<column name>)

FROM <table name>

WHERE <where predicates / conditions>

GROUP BY <list of columns>

HAVING <conditions>

Applied to the original data

Applies to the groups

Example:

```
SELECT      continent, AVG(LifeExpectancy) AS avg_le
FROM        country
WHERE       Population > 30000
GROUP BY    Continent
HAVING      avg_le > 70
```

Joins

SQL

Joins

Students Table

Student ID	First Name	Last Name	Major	GPA
125	Janet	Logan	EE	3.7
423	Janet	Carroll	CS	3.4
854	Farzad	Kamalzadeh	OR	3.6
239	Alex	Hagen	CE	3.9
371	Janet	Logan	EE	3.8

Courses Table

Course ID	Course Name
6458	Databases
7524	Big Data
6532	Python
4582	ML
3467	Data Mining

Enrollment Table

Course ID	Student ID
6458	423
7524	239
6532	125
4582	371
3467	854

Join of all the above tables

Student ID	First Name	Last Name	Major	GPA	Course ID	Course Name
125	Janet	Logan	EE	3.7	6532	Python
423	Janet	Carroll	CS	3.4	6458	Databases
854	Farzad	Kamalzadeh	OR	3.6	3467	Data Mining
239	Alex	Hagen	CE	3.9	7524	Big Data
371	Janet	Logan	EE	3.8	4582	ML

SQL


Join Types

Used to combine rows from two or more tables, based on a related column between them



(INNER)
JOIN

Returns records
that have matching
values in **both**
tables



LEFT
(OUTER)
JOIN

Returns all records
from the **left** table,
and the matched
records from the
right table



RIGHT
(OUTER)
JOIN

Returns all records
from the **right** table,
and the matched
records from the
left table



FULL
(OUTER)
JOIN

Returns all records
when there is a
match in **either left
or right** table

SQL

Joins

(INNER) JOIN



```
SELECT      <column list>
```

```
FROM        <table 1> AS T1, <table 2> AS T2
```

```
WHERE       T1.<col 1> = T2.<col 2>
```

Without the condition, it would be Cartesian JOIN, returning all possible combinations

Example:

```
SELECT      *  
FROM        country, city  
WHERE       country.Code = city.CountryCode
```

SQL

Joins

(INNER) JOIN



```
SELECT      <column list>
```

```
FROM        <table 1> AS T1
```

```
JOIN        <table 2> AS T2  
ON          T1.<col 1> = T2.<col 2>
```

Example:

```
SELECT      *  
FROM        customers AS c  
JOIN        sales AS s  
ON          c.customer_id = s.customer_id
```

SQL

Joins

LEFT (OUTER) JOIN



```
SELECT      <column list>
```

```
FROM        <table 1> AS T1
```

```
LEFT JOIN   <table 2> AS T2  
ON          T1.<col 1> = T2.<col 2>
```

Returns all records from the left table (table 1), and the matching records from the right table (table 2)

Example:

```
SELECT      *  
FROM        customers AS c  
LEFT JOIN   sales AS s  
ON          c.customer_id = s.customer_id
```

SQL

Joins

RIGHT (OUTER) JOIN



```
SELECT      <column list>
```

```
FROM        <table 1> AS T1
```

```
RIGHT JOIN  <table 2> AS T2  
ON          T1.<col 1> = T2.<col 2>
```

Returns all records from the right table (table 2), and the matching records from the left table (table 1)

Example:

```
SELECT      *  
FROM        customers AS c  
RIGHT JOIN  sales AS s  
ON          c.customer_id = s.customer_id
```

SQL

Joins

FULL (OUTER) JOIN



```
SELECT      <column list>
```

```
FROM        <table 1> AS T1
```

```
CROSS JOIN  <table 2> AS T2  
ON          T1.<col 1> = T2.<col 2>
```

Example:

```
SELECT      *  
FROM        customers AS c  
CROSS JOIN  sales AS s  
ON          c.customer_id = s.customer_id
```

SQL

Exercise

1. Exercise1

Table: Person

Column Name	Type
personId	int
lastName	varchar
firstName	varchar

personId is the primary key column for this table.

This table contains information about the ID of some persons and their first and last names.

Table: Address

Column Name	Type
addressId	int
personId	int
city	varchar
state	varchar

addressId is the primary key column for this table.

Each row of this table contains information about the city and state of one person with ID = PersonId.

Person table:

personId	lastName	firstName
1	Wang	Allen
2	Alice	Bob

Address table:

addressId	personId	city	state
1	2	New York City	New York
2	3	Apple	California

Output:

firstName	lastName	city	state
Allen	Wang	Null	Null
Bob	Alice	New York City	New York

Write an SQL query to report the first name, last name, city, and state of each person in the Person table. If the address of a personId is not present in the Address table, report null instead.

Return the result table in any order.

SQL

Answer

Person table:

personId	lastName	firstName
1	Wang	Allen
2	Alice	Bob

Address table:

addressId	personId	city	state
1	2	New York City	New York
2	3	Leetcode	California

```
SELECT      P.firstname, P.lastname, A.city, A.state
FROM        Person AS P
LEFT JOIN   Address AS A
ON          P.personId = A.personId
```

Output:

firstName	lastName	city	state
Allen	Wang	Null	Null
Bob	Alice	New York City	New York

SQL

Exercise

2. Exercise2

Table: Employee

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| id          | int    |
| name        | varchar|
| salary      | int    |
| managerId   | int    |
+-----+-----+
id is the primary key
column for this table.
Each row of this table
indicates the ID of an
employee, their name,
salary, and the ID of
their manager.
```

Write an SQL query to find the employees who earn more than their managers.
Return the result table in **any order**.

Employee table:

```
+----+-----+-----+-----+
| id | name  | salary | managerId |
+----+-----+-----+-----+
| 1  | Joe   | 70000  | 3         |
| 2  | Henry | 80000  | 4         |
| 3  | Sam   | 60000  | Null      |
| 4  | Max   | 90000  | Null      |
+----+-----+-----+-----+
```

Output:

```
+-----+
| Employee |
+-----+
| Joe      |
+-----+
```

SQL

Answer

Employee table:

id	name	salary	managerId
1	Joe	70000	3
2	Henry	80000	4
3	Sam	60000	Null
4	Max	90000	Null

```
SELECT      a.Name AS 'Employee'
FROM        Employee AS a,
            Employee AS b
WHERE       a.ManagerId = b.Id
AND         a.Salary > b.Salary
```

```
SELECT      a.NAME AS Employee
FROM        Employee AS a
JOIN        Employee AS b
ON          a.ManagerId = b.Id
AND         a.Salary > b.Salary
```

Output:

Employee
Joe

SQL

Exercise

3. Exercise3

Table: Customers

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| id          | int    |
| name        | varchar|
+-----+-----+
id is the primary key
column for this table.
Each row of this table
indicates the ID and name
of a customer.
```

Table: Orders

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| id          | int    |
| customerId  | int    |
+-----+-----+
id is the primary key
column for this table.
customerId is a
foreign key of the ID
from the Customers
table.
Each row of this table
indicates the ID of an
order and the ID of
the customer who
ordered it.
```

Customers table:

```
+-----+-----+
| id | name |
+-----+-----+
| 1  | Joe  |
| 2  | Henry|
| 3  | Sam  |
| 4  | Max  |
+-----+-----+
```

Orders table:

```
+-----+-----+
| id | customerId |
+-----+-----+
| 1  | 3          |
| 2  | 1          |
+-----+-----+
```

Output:

```
+-----+
| Customers |
+-----+
| Henry      |
| Max        |
+-----+
```

Write an SQL query to find the customer who did not order.
Return the result table in **any order**.

SQL

Answer

Customers table:

id	name
1	Joe
2	Henry
3	Sam
4	Max

Orders table:

id	customerId
1	3
2	1

```
SELECT      Name AS 'Customers'
FROM        Customers c
LEFT JOIN   Orders o
ON          c.Id = o.CustomerId
WHERE       o.CustomerId IS NULL
```

Output:

Customers
Henry
Max

SQL

Exercise

4. Exercise4

Table: Person

```
+-----+-----+
| Column Name | Type      |
+-----+-----+
| id          | int       |
| email       | varchar   |
+-----+-----+
```

id is the primary key column for this table.

Each row of this table contains an email. The emails will not contain uppercase letters.

Write an SQL query to report all the duplicate emails.

Return the result table in any order.

Person table:

```
+-----+-----+
| id | email |
+-----+-----+
| 1  | a@b.com |
| 2  | c@d.com |
| 3  | a@b.com |
+-----+-----+
```

Output:

```
+-----+
| Email |
+-----+
| a@b.com |
+-----+
```

SQL

Answer

Person table:

id	email
1	a@b.com
2	c@d.com
3	a@b.com

```
SELECT      Email
FROM        Person
GROUP BY    Email
HAVING      count(Email) > 1
```

Output:

Email
a@b.com

Sub-Queries

SQL

Sub-Queries

A Subquery or Inner query or a Nested query is a query within another SQL query and embedded within the SELECT clause (column Expression)



```
SELECT      <list of columns>, (sub-query)
```

```
FROM        <table name>
```

Example:

```
SELECT      emp_id, salary  
            (SELECT AVG(salary) AS avg_sal)  
            FROM employees  
FROM        employees
```



SQL

Sub-Queries

A Subquery or Inner query or a Nested query is a query within another SQL query and embedded within the FROM clause (table expression)



```
SELECT      <list of columns>
FROM        (Sub-Query) AS T1
```



Example:

```
SELECT      *
FROM        (SELECT      *
              FROM        country
              WHERE continent = 'Asia') AS T1
LIMIT      10
```

SQL

Sub-Queries

A Subquery or Inner query or a Nested query is a query within another SQL query and embedded within the WHERE clause



```
SELECT      <list of columns>
```

```
FROM        <table name>
```

```
WHERE       <condition with sub-query>
```

Example:

```
SELECT      *  
FROM        country  
WHERE       continent IN (SELECT continent  
                           FROM    country  
                           GROUP BY continent)
```

SQL

Exercise

5. Exercise5

Table: Customers

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| id          | int    |
| name        | varchar|
+-----+-----+
id is the primary key
column for this table.
Each row of this table
indicates the ID and name
of a customer.
```

Table: Orders

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| id          | int    |
| customerId  | int    |
+-----+-----+
id is the primary key
column for this table.
customerId is a
foreign key of the ID
from the Customers
table.
Each row of this table
indicates the ID of an
order and the ID of
the customer who
ordered it.
```

Customers table:

```
+-----+-----+
| id | name |
+-----+-----+
| 1  | Joe  |
| 2  | Henry|
| 3  | Sam  |
| 4  | Max  |
+-----+-----+
```

Orders table:

```
+-----+-----+
| id | customerId |
+-----+-----+
| 1  | 3          |
| 2  | 1          |
+-----+-----+
```

Output:

```
+-----+
| Customers |
+-----+
| Henry     |
| Max       |
+-----+
```

Write an SQL query to find the customer who did not order.
Return the result table in **any order**.

SQL

Answer

Customers table:

id	name
1	Joe
2	Henry
3	Sam
4	Max

Orders table:

id	customerId
1	3
2	1

```
SELECT      customers.name as 'Customers'
FROM        customers
WHERE       customers.id NOT IN
            (SELECT      customerId
             FROM        orders)
```

Output:

Customers
Henry
Max

SQL

Exercise

6. Exercise6

Table: Person

Column Name	Type
id	int
email	varchar

id is the primary key column for this table.

Each row of this table contains an email. The emails will not contain uppercase letters.

Person table:

id	email
1	a@b.com
2	c@d.com
3	a@b.com

Output:

Email
a@b.com

Write an SQL query to report all the duplicate emails.

Return the result table in any order.

SQL

Answer

Person table:

id	email
1	a@b.com
2	c@d.com
3	a@b.com

```
SELECT      Email
FROM        (SELECT Email, count(Email) as num
             FROM Person
             GROUP BY Email) AS T1
WHERE       num > 1
```

Output:

Email
a@b.com

SQL

More with Sub-Queries

To check sub-queries



EXISTS

To check if sub-query result is empty or not

ANY

To check if condition holds for some sub-query rows

ALL

To check if condition holds for all sub-query rows

Syntax:

EXISTS

(`<sub_query>`)

TRUE if the sub-query is non-empty

`<value> >= ANY (<sub_query>)`

TRUE if satisfied for some rows

`<value> >= ALL (<sub_query>)`


TRUE if satisfied for all rows

Set Operators

SQL

Set Operations


Used to combine, intersect or subtract rows from two or more tables




UNION



UNION ALL



INTERSECT



EXCEPT
(NOT IN)

NOTE: the queries must be compatible (have the same columns)

SQL

Set Operations

UNION (ALL)

Query 1 Result

*	*	*	*	*



Query 2 Result

*	*	*	*	*
*	*	*	*	*



UNION Result

*	*	*	*	*

UNION ALL Result

*	*	*	*	*
*	*	*	*	*

SQL

Set Operations

INTERSECT

Query 1 Result

*	*	*	*	*
+	+	+	+	+



Query 2 Result

*	*	*	*	*
+	+	+	+	+



INTERSECT
Result

*	*	*	*	*
+	+	+	+	+

SQL

Set Operations

EXCEPT (NOT IN)

Query 1 Result

+	+	+	+	+
#	#	#	#	#
*	*	*	*	*
-	-	-	-	-

—

Query 2 Result

+	+	+	+	+
-	-	-	-	-
*	*	*	*	*

=

EXCEPT Result

#	#	#	#	#

SQL

Set Operations

UNION (ALL), INTERSECT, EXCEPT



```
<query 1>
```

```
UNION/INTERSECT/EXCEPT
```

```
<query 2>
```

Example:

```
SELECT      a,b      FROM table1
UNION
SELECT      a,b      FROM table2
```

DDL

SQL

DDL



DDL
Data Definition Language



Define admissible database
content (schema)



DDL
(Data
Definition
Language)

CREATE
DROP
ALTER
RENAME

Relations and their Schemata

What columns each table have and what are the
column types

Constraints to restrict admissible contents

Constraints on single relations and constraints linking
multiple relations

SQL

DDL: Table Creation



Table Creation

To create a table, we need to know the following:

Schema

Table
Name

Column
Names

Data
Types

Allow
Duplicates

Allow
Nulls

Constraints

SQL

DDL: Table Creation

Table Creation Syntax



```
CREATE TABLE <table name>
            ( <table definition> )
```

<table definition> is a comma-separated column definition: <column name> <column type>

Example:

```
CREATE TABLE students
            (student_id    INT,
             first_name    VARCHAR(20) ,
             last_name     VARCHAR(20) )
```

SQL

DDL: Data Types



Data types (feature/column types)



Character String

Fixed length: `CHAR()`

Variable Length: `VARCHAR()`



Numeric

Integer: `INT()`, `SMALLINT`, `BIGINT`

Decimal:
`DECIMAL(n, m)`
`DOUBLE`
`FLOAT`
`REAL(n, m)`



Boolean

True/False: `BOOLEAN`



Date/Time

Date: `DATE`

Time: `TIME()`

Date & Time: `DATETIME()`

Time Stamp: `TIMESTAMP()`

SQL

DDL: Data Types



Character String

Fixed length: `CHAR ()`

Variable Length: `VARCHAR ()`

Value	CHAR(4)	Storage Required	VARCHAR(4)	Storage Required
"	' '	4 bytes	"	1 byte
'ab'	'ab '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefgh'	'abcd'	4 bytes	'abcd'	5 bytes

SQL

DDL: Data Types

123

Numeric

Integer: `INT()` , `SMALLINT` , `BIGINT`

Type	Storage (Bytes)	Minimum Value Signed	Minimum Value Unsigned	Maximum Value Signed	Maximum Value Unsigned
TINYINT	1	-128	0	127	255
SMALLINT	2	-32768	0	32767	65535
MEDIUMINT	3	-8388608	0	8388607	16777215
INT	4	-2147483648	0	2147483647	4294967295
BIGINT	8	-2^{63}	0	$2^{63}-1$	$2^{64}-1$

SQL

DDL: Data Types

123

Numeric

Decimal:

DECIMAL (n , m)
DOUBLE
FLOAT
REAL (n , m)

Fixed-Point Types
(Exact Value)
- DECIMAL, NUMERIC

Floating-Point Types
(Approximate Value)
- FLOAT, DOUBLE

SQL

DDL: Data Types

Numeric:

Fixed-Point Types (Exact Value) - DECIMAL, NUMERIC

- The DECIMAL and NUMERIC types store exact numeric data values. These types are used when it is important to preserve exact precision, for example with monetary data.
- The maximum number of digits for DECIMAL is 65, but the actual range for a given DECIMAL column can be constrained by the precision or scale for a given column. When such a column is assigned a value with more digits following the decimal point than are permitted by the specified scale, the value is converted to that scale. (The precise behavior is operating system-specific, but generally the effect is truncation to the permissible number of digits.)

Floating-Point Types (Approximate Value) - FLOAT, DOUBLE

- The FLOAT and DOUBLE types represent approximate numeric data values. MySQL uses four bytes for single-precision values and eight bytes for double-precision values.
- Because floating-point values are approximate and not stored as exact values, attempts to treat them as exact in comparisons may lead to problems. They are also subject to platform or implementation dependencies.
- <https://stackoverflow.com/questions/1056323/difference-between-numeric-float-and-decimal-in-sql-server>

SQL

DDL: Data Types



Date/Time

Date: **DATE**

Time: **TIME (fsp)**

Date & Time: **DATETIME (fsp)**

Time Stamp: **TIMESTAMP (fsp)**

The **fsp** value, if given, must be in the range 0 to 6. A value of 0 signifies that there is no fractional part. If omitted, the default precision is 0.

```
CREATE TABLE t1
    (t TIME (3) ,
     dt DATETIME (6) ,
     ts TIMESTAMP (0) ) ;
```

Data Type	Example	Format
<u>DATE</u>	'1992-02-10'	'YYYY-MM-DD'
<u>TIME</u>	'10:35:24'	<i>hh:mm:ss[.fraction]</i>
<u>DATETIME</u>	'1992-02-10 10:35:24'	'YYYY-MM-DD hh:mm:ss[.fraction]'
<u>TIMESTAMP</u>	'1992-02-10 10:35:24'	'YYYY-MM-DD hh:mm:ss[.fraction]'
<u>YEAR</u>	1992	YYYY

SQL

DDL: Data Types



Date/Time

Date: **DATE**

Time: **TIME ()**

Date & Time: **DATETIME ()**

Time Stamp: **TIMESTAMP ()**

- The SUM() and AVG() aggregate functions do not work with temporal values.
 - (They convert the values to numbers, losing everything after the first nonnumeric character.)
- To work around this problem, convert to numeric units, perform the aggregate operation, and convert back to a temporal value.



```
SELECT      SEC_TO_TIME ( SUM(    TIME_TO_SEC ( time_col)    ) )  
FROM  
tbl_name;
```

```
SELECT      FROM_DAYS (    SUM(    TO_DAYS (      date_col)    ) )  
FROM  
tbl_name;
```


SQL

DDL: Data Types



CAST Function

The CAST() function in MySQL is used to convert a value from one data type to another data type specified in the expression.



CAST(<expression> AS <datatype>)

- Expression: It is a value that will be converted into another specific datatype.
- Datatype: It is a value or data type in which the expression value needs to be converted.

Example:

```
SELECT      CAST("2018-11-30" AS DATE);  
SELECT      CAST(3-6 AS SIGNED);  
SELECT      CONCAT('CAST Function Example ## ', CAST(5 AS CHAR));
```

SQL

DDL: Data Types



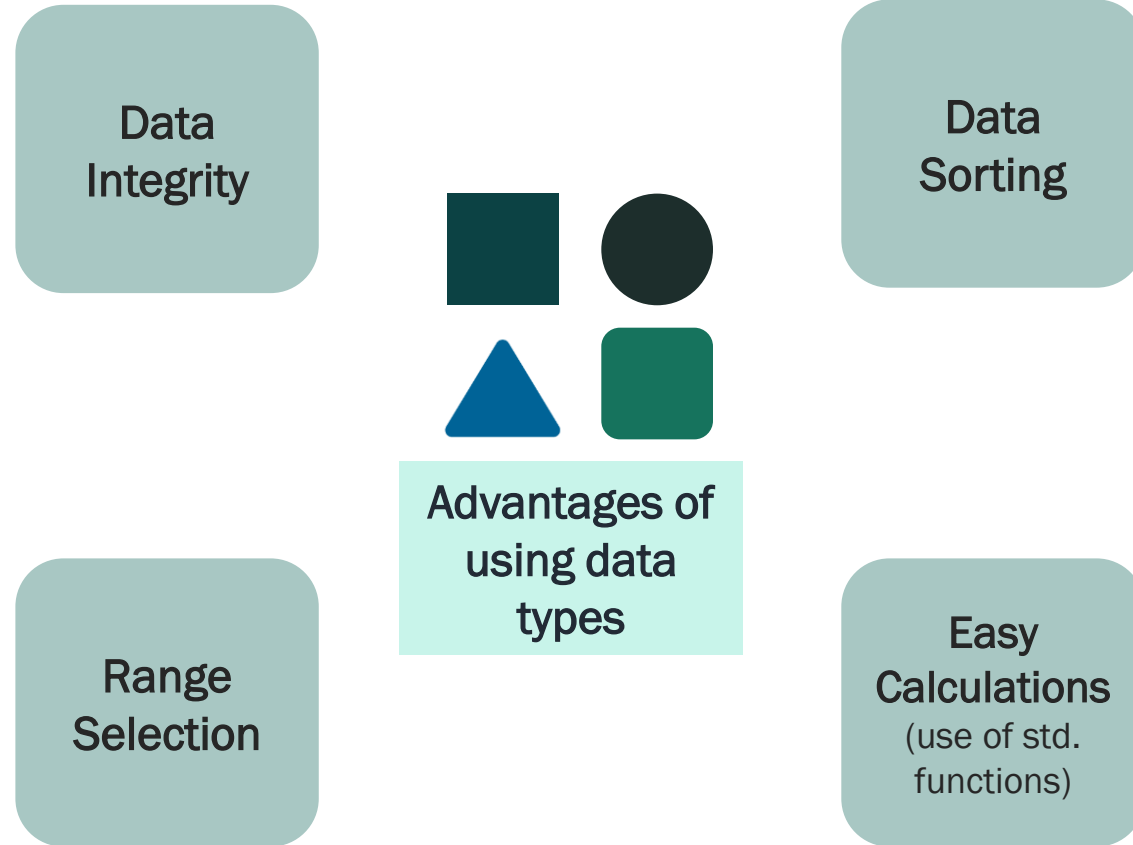
Date/Time Functions

Below is a list of useful date and time functions in MySQL
(A complete [List](#) of these functions)

Name	Description
CURDATE()	Return the current date
CURTIME()	Return the current time
DATE_FORMAT()	Format date as specified
DATEDIFF()	Subtract two dates
NOW()	Return the current date and time
STR_TO_DATE()	Convert a string to a date
TIME_FORMAT()	Format as time
TIMEDIFF()	Subtract time
TIMESTAMP()	With a single argument, this function returns the date or datetime expression; with two arguments, the sum of the arguments
TIMESTAMPDIFF()	Subtract an interval from a datetime expression

SQL

DDL: Data Types



SQL

DDL: Primary Key Constraint

Table Creation Syntax with Primary Key



```
CREATE TABLE <table name>
    ( <table definition>,
      PRIMARY KEY (<column name>) )
```

```
CREATE TABLE <table name>
    ( <table definition>,
      CONSTRAINT <constraint name>
      PRIMARY KEY (<column names>) )
```

Example:

```
CREATE TABLE students
    (student_id    INT,
     first_name    VARCHAR(20),
     last_name     VARCHAR(20),
     PRIMARY KEY  (student_id))
```

SQL

DDL: Foreign Key Constraint

Table Creation Syntax with Foreign Key



```
CREATE TABLE <table name>
    ( <table definition>,
      CONSTRAINT <constraint name>
      FOREIGN KEY (<column name>, ...)
      REFERENCES <tbl_name> (<col_name>, ...)
      ON DELETE <reference_option>
      ON UPDATE <reference_option>)
```

reference_option:

RESTRICT | CASCADE | SET NULL | NO ACTION | SET DEFAULT

Example:

```
CREATE TABLE      enrollment
    (student_id      INT,
     first_name      VARCHAR(20),
     last_name       VARCHAR(20),
     CONSTRAINT fk1
        FOREIGN KEY (student_id)
        REFERENCES students(std_id)
        ON DELETE CASCADE
        ON UPDATE NO ACTION)
```

- **Referential Actions**

- When an UPDATE or DELETE operation affects a key value in the parent table that has matching rows in the child table, the result depends on the referential action specified by ON UPDATE and ON DELETE subclauses of the FOREIGN KEY clause. Referential actions include:
- **CASCADE:** Delete or update the row from the parent table and automatically delete or update the matching rows in the child table. Both ON DELETE CASCADE and ON UPDATE CASCADE are supported. Between two tables, do not define several ON UPDATE CASCADE clauses that act on the same column in the parent table or in the child table.
- If a FOREIGN KEY clause is defined on both tables in a foreign key relationship, making both tables a parent and child, an ON UPDATE CASCADE or ON DELETE CASCADE subclause defined for one FOREIGN KEY clause must be defined for the other in order for cascading operations to succeed. If an ON UPDATE CASCADE or ON DELETE CASCADE subclause is only defined for one FOREIGN KEY clause, cascading operations fail with an error.
- **SET NULL:** Delete or update the row from the parent table and set the foreign key column or columns in the child table to NULL. Both ON DELETE SET NULL and ON UPDATE SET NULL clauses are supported.
- If you specify a SET NULL action, make sure that you have not declared the columns in the child table as NOT NULL.
- **RESTRICT:** Rejects the delete or update operation for the parent table. Specifying RESTRICT (or NO ACTION) is the same as omitting the ON DELETE or ON UPDATE clause.
- **NO ACTION:** A keyword from standard SQL. In MySQL, equivalent to RESTRICT. The MySQL Server rejects the delete or update operation for the parent table if there is a related foreign key value in the referenced table. Some database systems have deferred checks, and NO ACTION is a deferred check. In MySQL, foreign key constraints are checked immediately, so NO ACTION is the same as RESTRICT.
- **SET DEFAULT:** This action is recognized by the MySQL parser, but both InnoDB and NDB reject table definitions containing ON DELETE SET DEFAULT or ON UPDATE SET DEFAULT clauses.

SQL

DDL: Alteration



Alteration

Alter is used for:

Add/remove columns

Modify the data types

Add/remove keys

Add/remove constraints

SQL

DDL: Alteration

Add columns



```
ALTER TABLE      <table name>
ADD COLUMN        <column name> <data type>
                  <other options>
```

Other options:

```
NULL | NOT NULL | FIRST | AFTER <column name> | BINARY | ...
```

Example:

```
ALTER TABLE      students
ADD COLUMN        age INT NULL AFTER last_name
```


SQL

DDL: Alteration

Remove columns



```
ALTER TABLE      <table name>  
DROP COLUMN      <column name>
```

Example:

```
ALTER TABLE  students  
DROP COLUMN  age
```

SQL

DDL: Alteration

Change/modify/rename columns



```
ALTER TABLE      <table name>
CHANGE COLUMN      <column name>
                   <new column name> <new definition>
```

```
ALTER TABLE      <table name>
MODIFY             <column name> <new definition>
```

Example:

```
ALTER TABLE      students
CHANGE COLUMN      last_name
                   lname VARCHAR(20) NOT NULL
```

SQL

DDL: Alteration

Rename table



```
ALTER TABLE      <table name>  
RENAME TO        <new table name>
```

Example:

```
ALTER TABLE      students  
RENAME TO        students_table
```

SQL

DDL: Alteration

Add Primary Key constraint



```
ALTER TABLE      <table name>  
ADD PRIMARY KEY   (<column name>)
```

```
ALTER TABLE      <table name>  
ADD CONSTRAINT    <constraint name>  
PRIMARY KEY       (<column name>)
```

Example:

```
ALTER TABLE      students  
ADD PRIMARY KEY   (student_id)
```

SQL

DDL: Alteration

Drop Primary Key constraint



```
ALTER TABLE      <table name>  
DROP PRIMARY KEY
```

Example:

```
ALTER TABLE      students  
DROP PRIMARY KEY
```

SQL

DDL: Alteration

Add Foreign Key constraint



```
ALTER TABLE      <table name>
ADD FOREIGN KEY   (<column name>)
REFERENCES        <reference table>(<reference column>)
```

```
ALTER TABLE      <table name>
ADD CONSTRAINT    <constraint name>
FOREIGN KEY       (<column name>)
REFERENCES        <reference table>(<reference column>)
```

Example:

```
ALTER TABLE      enrollment
ADD CONSTRAINT    fk_enrol_std
FOREIGN KEY       (student_id)
REFERENCES        students(std_id)
```

DML

SQL

DML



DML

Data Manipulation Language



Used to manipulate the data of
an existing schema



DML

(Data
Manipulation
Language)

INSERT
UPDATE
DELETE


Adding/inserting data into a table

Deleting/removing data from a table

Updating existing data from a table

SQL

DML: Data Insertion



Data Insertion into a table

DML is used to insert data into an existing table

Add a complete row
(fully specified row)

Add a part of a row
(partially specified row)

Add data from a file

SQL

DML: Insertion

Fully specified row



```
INSERT INTO      <table name>
VALUES           (<list of values>)
```

```
INSERT INTO      <table name>
VALUES           (<value list 1>),
                  (<value list 2>),
                  (<value list 3>), ...
```

Values should be separated by comma and in the same order as the columns

Example:

```
INSERT INTO      students
VALUES           (765427, 'Alex', 'Hagen')
```

SQL

DML: Insertion

Partially specified row



```
INSERT INTO      <table name> (<list of columns>)  
VALUES          (<list of values>)
```

Values should be separated by comma and in the same order as the columns

Example:

```
INSERT INTO      students (std_id, first_name)  
VALUES          (765427, 'Alex')
```

SQL

DML: Insertion

From a file



```
LOAD DATA (LOCAL) INFILE    '<location of the file>'
INTO TABLE                  <table name>
FIELDS TERMINATED BY        '<field delimiter>'
ENCLOSED BY                  '<string identifier>'
LINES TERMINATED BY          '<new line identifier>'
IGNORE 1 ROWS;
```

First row can be ignored if it contains the header (names of the columns)

<https://dev.mysql.com/doc/refman/5.7/en/load-data.html>

Example:

```
LOAD DATA INFILE            'c:/temp/students.csv'
INTO TABLE                  students
FIELDS TERMINATED BY        ','
ENCLOSED BY                  '"'
LINES TERMINATED BY          '\n'
IGNORE 1 ROWS;
```

<https://stackoverflow.com/questions/32737478/how-should-i-resolve-secure-file-priv-in-mysql>

SQL

DML: Data Deletion



Deleting data from a table

DML is used to delete data from an existing table



```
DELETE FROM      <table name>  
WHERE            <condition>
```

Example:

```
DELETE FROM      students  
WHERE            std_id = 126417
```

SQL

DML: Data Update



Updating data in a table

DML is also used to update data in an existing table



```
UPDATE      <table name>
SET         <column name> = <value>
WHERE      <condition>
```

Example:

```
UPDATE      students
SET         last_name = 'Hagen'
WHERE      std_id = 765427
```

Indexes

SQL

Indexes

How to create an index



```
CREATE INDEX      <index name>  
ON                <table name> (<list of columns>)
```

List of columns is the column names separated by commas.

Example:

```
CREATE INDEX      emp_name_idx  
ON                employees (emp_first_name)
```

```
CREATE INDEX      emp_name_idx  
ON                employees (emp_first_name(10))
```


SQL

Indexes

How to see all the indexes from a table



```
SHOW INDEX  
FROM          <table name>  
FROM          <database name>
```

Lists all the indexes from a given table in a given database

Example:

```
SHOW INDEX  
FROM          students  
FROM          university_db
```

Transactions

SQL

Transactions

How to create a transaction



```
START TRANSACTION;
```

```
<query>;
```

```
<query>;
```

```
...
```

```
<query>;
```

```
COMMIT;
```

Options:

```
SET autocommit = 0;
```

Disables the autocommit option

```
ROLLBACK
```

Rolls back the changes made instead of committing them

Query Timing

SQL

Query Timing

How to time your queries



```
SET PROFILING = 1;
```

```
<query 1>;
```

```
<query 2>;
```

```
...
```

```
<query n>;
```


```
SHOW PROFILES;
```

Show profile will return a table of the queries executed and the time it took for them to run

SQL

MQL

Basic Commands



show dbs

Show all the databases

use <db name>

Switches to the given db, creates one if it does not exist

db

Shows the current db being worked on


show collections

Shows all the collections in the current db

<https://www.mongodb.com/developer/quickstart/cheat-sheet/>

MQL

Query



```
db.<collection name>.find()
```

Shows all the documents in the collection

```
db.<collection name>.find().pretty()
```

Show all the docs in the coll, formatted

```
db.<collection name>.findOne()
```


Returns a single document

```
db.<collection name>.distinct("<property>")
```

Returns all the distinct values the given property has in the database

MQL

Query



```
db.<collection name>.find({specify a condition})
```

Shows all the matches from the database

```
db.<collection name>.find({$or: [{condition1}, {condition2}]})
```

Shows all the matches from the database given the list of conditions

Example:

```
db.students.find({name: "Max"})
```


```
db.students.find({date: ISODate("2020-09-25")})
```

```
db.students.find({$and: [{name: "Max"}, {major: "EE"}]})
```

```
db.students.find({name: "Max", age: 32})
```

MQL

Query - Counts



```
db.<collection name>.count()
```

Counts the number of documents that match the criteria (based on collection metadata)

```
db.<collection name>.countDocuments()
```

Counts the number of documents that match the criteria (accurate count)

Example:

```
db.students.count({age: 32})
```

```
db.students.countDocuments({age: 32})
```

MQL

Comparison Operators

Comparison Operators



`$eq`

`$ne`

`$gt`

`$lt`

`$gte`

`$lte`

`$in`

`$nin`

Used to compare the values of two operands

Example:

```
db.students.find({"year": {$gt: 1970}})
db.students.find({"year": {$gte: 1970}})
db.students.find({"year": {$lt: 1970}})
db.students.find({"year": {$lte: 1970}})
db.students.find({"year": {$ne: 1970}})
db.students.find({"year": {$in: [1958, 1959]}})
db.students.find({"year": {$nin: [1958, 1959]}})
```

MQL

Logical Operators

Logical Operators



\$not

\$or

\$nor

\$and

Used to combine logical operations

Example:

```
db.students.find({name:{$not: {$eq: "Max"}}})
db.students.find({$or: [{"year" : 1958}, {"year" : 1959}]})
db.students.find({$nor: [{gpa: 1.99}, {failed: true}]})
db.students.find({
  $and: [
    {$or: [{gpa: {$lt: 3}}, {gpa :{$gt: 2}}]},
    {$or: [{failed: true}, {gpa: {$lt: 3 }}]}
  ]
})
```

MQL

Element Operators

Element Operators



\$exists

Used to check if a specific property exists

\$type

Used for checking the value types in the database


<https://www.mongodb.com/docs/manual/reference/operator/query/type/>

Example:

```
db.students.find({name: {$exists: true}})
db.students.find({"zipCode": {$type: 2 }})
db.students.find({"zipCode": {$type: "string"}})
```

MQL

Create



```
db.createCollection('<collection name>')
```

Creates a collection

```
db.<collection name>.insertOne()
```

Inserts a new document into the collection


Example:

```
db.createCollection('students')
```

```
db.students.insertOne({name: "Max"})
```

MQL

Create



```
db.<collection name>.insert()
```

Inserts a new document into the collection

```
db.<collection name>.insert([,])
```

Inserts multiple documents into the collection

Example:


```
db.students.insert([ {name: "Max"}, {name: "Alex"} ])
```

```
db.students.insert({date: new Date('2021-11-21')})
```

<https://www.mongodb.com/docs/manual/reference/method/Date/>

MQL

Update



```
db.<collection name>.update({<condition>},{update})
```

replaces an existing document with the new fields and values

```
db.<collection name>.update({<condition>}, {$set:{update}})
```

Updates an existing document keeping the rest of the properties untouched

<https://www.mongodb.com/docs/manual/reference/method/db.collection.update/>


Example:

```
db.students.update({"_id": 1}, {"year": 2016})
```

```
db.students.update({"_id": 1}, {$set: {"year": 2016, name: "Max" }})
```


MQL

Update



```
db.<collection name>.update({<condition>}, {$unset: {update}})
```

Removes a property

```
db.<collection name>.update({<condition>}, {$rename: {"<f>": "<g>"}})
```

Renames field (property) f to g


Example:

```
db.students.update({"_id": 1}, {$unset: {"year": 1}})
```

```
db.students.update({"_id": 1}, {$rename: {"year": "date"} })
```

MQL

Delete



```
db.<collection name>.remove({<condition>})
```

Removes documents based on the condition

```
db.<collection name>.findOneAndDelete({<condition>})
```

Remove one document based on the condition

Example:

```
db.students.remove({name: "Max"})
```

```
db.students.findOneAndDelete({"name": "Max"})
```



<https://www.mongodb.com/docs/manual/tutorial/query-embedded-documents/>

Exercises

SQL

Exercise

197. Rising Temperature

Table: Weather

Column Name	Type
id	int
recordDate	date
temperature	int

id is the primary key for this table.

This table contains information about the temperature on a certain day.

Write an SQL query to find all dates' Id with higher temperatures compared to its previous dates (yesterday).

Return the result table in any order.

Weather table:

id	recordDate	temperature
1	2015-01-01	10
2	2015-01-02	25
3	2015-01-03	20
4	2015-01-04	30

Output:

id
2
4

SQL

Answer

Weather table:

id	recordDate	temperature
1	2015-01-01	10
2	2015-01-02	25
3	2015-01-03	20
4	2015-01-04	30

```
SELECT      weather.id AS 'Id'
FROM        weather
JOIN        weather w
ON          DATEDIFF(weather.recordDate, w.recordDate) = 1
AND         weather.Temperature > w.Temperature;
```

Output:

id
2
4

SQL

Exercise

603. Consecutive Available Seats

Table: Cinema

```
+-----+-----+
| Column Name | Type |
+-----+-----+
| seat_id     | int  |
| free        | bool |
+-----+-----+
seat_id is an auto-
increment primary key
column for this table.
Each row of this table
indicates whether the
ith seat is free or
not. 1 means free
while 0 means
occupied.
```

Write an SQL query to report all the consecutive available seats in the cinema.

Return the result table ordered by seat_id in ascending order.

Cinema table:

```
+-----+-----+
| seat_id | free |
+-----+-----+
| 1       | 1    |
| 2       | 0    |
| 3       | 1    |
| 4       | 1    |
| 5       | 1    |
+-----+-----+
```

Output:

```
+-----+
| seat_id |
+-----+
| 3       |
| 4       |
| 5       |
+-----+
```

SQL

Answer

```
SELECT      a.seat_id, a.free, b.seat_id, b.free
FROM        cinema a
JOIN        cinema b;
```

```
SELECT      a.seat_id, a.free, b.seat_id, b.free
FROM        cinema a
JOIN        cinema b
ON          ABS(a.seat_id - b.seat_id) = 1
AND a.free = true and b.free = true;
```

```
SELECT      distinct a.seat_id
FROM        cinema a
JOIN        cinema b
ON          ABS(a.seat_id - b.seat_id) = 1
AND a.free = true and b.free = true
ORDER BY    a.seat_id;
```

Cinema table:

seat_id	free
1	1
2	0
3	1
4	1
5	1

Output:

seat_id
3
4
5

SQL

Exercise

610. Triangle Judgement

Table: Triangle

```
+-----+-----+
| Column Name | Type |
+-----+-----+
| x           | int  |
| y           | int  |
| z           | int  |
+-----+-----+
```

(x, y, z) is the primary key column for this table.

Each row of this table contains the lengths of three line segments.

Write an SQL query to report for every three line segments whether they can form a triangle.

Return the result table in any order.

Triangle table:

```
+-----+-----+-----+
| x  | y  | z  |
+-----+-----+-----+
| 13 | 15 | 30 |
| 10 | 20 | 15 |
+-----+-----+-----+
```

Output:

```
+-----+-----+-----+-----+
| x  | y  | z  | triangle |
+-----+-----+-----+-----+
| 13 | 15 | 30 | No        |
| 10 | 20 | 15 | Yes       |
+-----+-----+-----+-----+
```

SQL

Answer

Triangle table:

x	y	z
13	15	30
10	20	15

Output:

x	y	z	triangle
13	15	30	No
10	20	15	Yes

```
SELECT      x, y, z,
            CASE
                WHEN x + y > z AND x + z > y AND y + z > x
                    THEN 'Yes'
                ELSE 'No'
            END AS 'triangle'
FROM        triangle;
```

SQL

Exercise

627. Swap Salary

Table: Salary

Column Name	Type
id	int
name	varchar
sex	ENUM
salary	int

id is the primary key for this table.

The sex column is ENUM value of type ('m', 'f').

The table contains information about an employee.

Write an SQL query to swap all 'f' and 'm' values (i.e., change all 'f' values to 'm' and vice versa) with a single update statement and no intermediate temporary tables.

Note that you must write a single update statement, do not write any select statement for this problem.

Salary table:

id	name	sex	salary
1	A	m	2500
2	B	f	1500
3	C	m	5500
4	D	f	500

Output:

id	name	sex	salary
1	A	f	2500
2	B	m	1500
3	C	f	5500
4	D	m	500

SQL

Answer

Salary table:

id	name	sex	salary
1	A	m	2500
2	B	f	1500
3	C	m	5500
4	D	f	500

Output:

id	name	sex	salary
1	A	f	2500
2	B	m	1500
3	C	f	5500
4	D	m	500

```
UPDATE salary
SET sex = CASE sex
            WHEN 'm'
            THEN 'f'
            ELSE 'm'
            END;
```

SQL

Exercise

1075. Project Employees I

Table: Project

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| project_id  | int    |
| employee_id | int    |
+-----+-----+
(project_id, employee_id)
is the primary key of
this table.
employee_id is a foreign
key to Employee table.
Each row of this table
indicates that the
employee with employee_id
is working on the project
with project_id.
```

Table: Employee

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| employee_id | int    |
| name        | varchar|
| experience_years | int    |
+-----+-----+
employee_id is the primary key
of this table.
Each row of this table contains
information about one employee.
```

Project table:

```
+-----+-----+
| project_id | employee_id |
+-----+-----+
| 1          | 1           |
| 1          | 2           |
| 1          | 3           |
| 2          | 1           |
| 2          | 4           |
+-----+-----+
```

Employee table:

```
+-----+-----+-----+
| employee_id | name   | experience_years |
+-----+-----+-----+
| 1           | Khaled | 3                |
| 2           | Ali    | 2                |
| 3           | John   | 1                |
| 4           | Doe    | 2                |
+-----+-----+-----+
```

Write an SQL query that reports the average experience years of all the employees for each project, rounded to 2 digits.

Return the result table in any order.

```
+-----+-----+
| project_id | average_years |
+-----+-----+
| 1          | 2.00          |
| 2          | 2.50          |
+-----+-----+
```

SQL

Answer

Project table:

project_id	employee_id
1	1
1	2
1	3
2	1
2	4

Employee table:

employee_id	name	experience_years
1	Khaled	3
2	Ali	2
3	John	1
4	Doe	2

project_id	average_years
1	2.00
2	2.50

```
SELECT      p.project_id,
            ROUND(AVG(e.experience_years), 2) AS average_years
FROM        project p
JOIN        employee e
ON          p.employee_id = e.employee_id
GROUP BY    1;
```

SQL

Exercise

1076. Project Employees II

Table: Project

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| project_id  | int    |
| employee_id | int    |
+-----+-----+
(project_id, employee_id)
is the primary key of
this table.
employee_id is a foreign
key to Employee table.
Each row of this table
indicates that the
employee with employee_id
is working on the project
with project_id.
```

Table: Employee

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| employee_id  | int    |
| name         | varchar|
| experience_years | int    |
+-----+-----+
employee_id is the primary key
of this table.
Each row of this table contains
information about one employee.
```

Project table:

```
+-----+-----+
| project_id | employee_id |
+-----+-----+
| 1          | 1           |
| 1          | 2           |
| 1          | 3           |
| 2          | 1           |
| 2          | 4           |
+-----+-----+
```

Employee table:

```
+-----+-----+-----+
| employee_id | name   | experience_years |
+-----+-----+-----+
| 1          | Khaled | 3                |
| 2          | Ali    | 2                |
| 3          | John   | 1                |
| 4          | Doe    | 2                |
+-----+-----+-----+
```

Output:

```
+-----+
| project_id |
+-----+
| 1          |
+-----+
```

Write an SQL query that reports all the projects that have the most employees.

Return the result table in any order.

SQL

Answer

Project table:

project_id	employee_id
1	1
1	2
1	3
2	1
2	4

Employee table:

employee_id	name	experience_years
1	Khaled	3
2	Ali	2
3	John	1
4	Doe	2

Output:

project_id
1

```
WITH employee_count AS (SELECT project_id, COUNT(employee_id) AS cnt
                        FROM Project
                        GROUP BY project_id)

SELECT project_id
FROM employee_count
WHERE cnt = (SELECT MAX(cnt)
            FROM employee_count);
```


SQL

Exercise

1082. Sales Analysis I

Table: Sales

Column Name	Type
seller_id	int
product_id	int
buyer_id	int
sale_date	date
quantity	int
price	int

This table has no primary key, it can have repeated rows.

product_id is a foreign key to the Product table. Each row of this table contains some information about one sale.

Write an SQL query that reports the best seller by total sales price, If there is a tie, report them all.

Return the result table in any order.

Table: Product

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product_id is the primary key of this table.

Each row of this table indicates the name and the price of each product.

Output:

seller_id
1
3

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

SQL

Answer

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

```
SELECT      seller_id
FROM        Sales
GROUP BY    seller_id
HAVING      SUM(price) = (SELECT SUM(price)
                           FROM Sales
                           GROUP BY seller_id
                           ORDER BY 1 DESC
                           LIMIT 1)
```

Output:

seller_id
1
3

SQL

Exercise

1083. Sales Analysis II

Table: Sales

Column Name	Type
seller_id	int
product_id	int
buyer_id	int
sale_date	date
quantity	int
price	int

This table has no primary key, it can have repeated rows.

product_id is a foreign key to the Product table. Each row of this table contains some information about one sale.

Write an SQL query that reports the buyers who have bought S8 but not iPhone. Note that S8 and iPhone are products present in the Product table.

Return the result table in any order.

Table: Product

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product_id is the primary key of this table.

Each row of this table indicates the name and the price of each product.

Output:

buyer_id
1

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

SQL

Answer

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

```
SELECT      b.buyer_id
FROM        Product AS a
JOIN        Sales AS b
ON          a.product_id = b.product_id
GROUP BY    b.buyer_id
HAVING      SUM(a.product_name = 'S8') > 0
            AND SUM(a.product_name = 'iPhone') = 0;
```

Output:

buyer_id
1

SQL

Exercise

1084. Sales Analysis III

Table: Sales

Column Name	Type
seller_id	int
product_id	int
buyer_id	int
sale_date	date
quantity	int
price	int

This table has no primary key, it can have repeated rows.

product_id is a foreign key to the Product table. Each row of this table contains some information about one sale.

Table: Product

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product_id is the primary key of this table.

Each row of this table indicates the name and the price of each product.

Output:

product_id	product_name
1	S8

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

Write an SQL query that reports the products that were only sold in the spring of 2019. That is, between 2019-01-01 and 2019-03-31 inclusive.

Return the result table in any order.

SQL

Answer

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

```
SELECT      s.product_id, product_name
FROM        Sales s
LEFT JOIN   Product p
ON          s.product_id = p.product_id
GROUP BY    s.product_id
HAVING      MIN(sale_date) >= CAST('2019-01-01' AS DATE)
AND         MAX(sale_date) <= CAST('2019-03-31' AS DATE)
```

Output:

product_id	product_name
1	S8

SQL

Exercise

1113. Reported Posts

Table: Actions

Column Name	Type
user_id	int
post_id	int
action_date	date
action	enum
extra	varchar

There is no primary key for this table, it may have duplicate rows.

The action column is an ENUM type of ('view', 'like', 'reaction', 'comment', 'report', 'share').

The extra column has optional information about the action, such as a reason for the report or a type of reaction.

Write an SQL query that reports the number of posts reported yesterday for each report reason. Assume today is 2019-07-05.

Return the result table in any order.

Input:

Actions table:

user_id	post_id	action_date	action	extra
1	1	2019-07-01	view	null
1	1	2019-07-01	like	null
1	1	2019-07-01	share	null
2	4	2019-07-04	view	null
2	4	2019-07-04	report	spam
3	4	2019-07-04	view	null
3	4	2019-07-04	report	spam
4	3	2019-07-02	view	null
4	3	2019-07-02	report	spam
5	2	2019-07-04	view	null
5	2	2019-07-04	report	racism
5	5	2019-07-04	view	null
5	5	2019-07-04	report	racism

Output:

report_reason	report_count
spam	1
racism	2

SQL

Answer

Input:

Actions table:

user_id	post_id	action_date	action	extra
1	1	2019-07-01	view	null
1	1	2019-07-01	like	null
1	1	2019-07-01	share	null
2	4	2019-07-04	view	null
2	4	2019-07-04	report	spam
3	4	2019-07-04	view	null
3	4	2019-07-04	report	spam
4	3	2019-07-02	view	null
4	3	2019-07-02	report	spam
5	2	2019-07-04	view	null
5	2	2019-07-04	report	racism
5	5	2019-07-04	view	null
5	5	2019-07-04	report	racism

Output:

report_reason	report_count
spam	1
racism	2

```
SELECT      extra AS report_reason, count(distinct post_id) AS report_count
FROM        actions
WHERE       action_date = '2019-07-04' AND action = 'report'
GROUP BY    extra
```


SQL

Exercise

1141. User Activity for the Past 30 Days I

Table: Activity

Column Name	Type
user_id	int
session_id	int
activity_date	date
activity_type	enum

There is no primary key for this table, it may have duplicate rows.

The activity_type column is an ENUM of type ('open_session', 'end_session', 'scroll_down', 'send_message').

The table shows the user activities for a social media website.

Note that each session belongs to exactly one user.

Write an SQL query to find the daily active user count for a period of 30 days ending 2019-07-27 inclusively. A user was active on someday if they made at least one activity on that day.

Return the result table in any order.

Input:

Activity table:

user_id	session_id	activity_date	activity_type
1	1	2019-07-20	open_session
1	1	2019-07-20	scroll_down
1	1	2019-07-20	end_session
2	4	2019-07-20	open_session
2	4	2019-07-21	send_message
2	4	2019-07-21	end_session
3	2	2019-07-21	open_session
3	2	2019-07-21	send_message
3	2	2019-07-21	end_session
4	3	2019-06-25	open_session
4	3	2019-06-25	end_session

Output:

day	active_users
2019-07-20	2
2019-07-21	2

SQL

Answer

Input:

Activity table:

user_id	session_id	activity_date	activity_type
1	1	2019-07-20	open_session
1	1	2019-07-20	scroll_down
1	1	2019-07-20	end_session
2	4	2019-07-20	open_session
2	4	2019-07-21	send_message
2	4	2019-07-21	end_session
3	2	2019-07-21	open_session
3	2	2019-07-21	send_message
3	2	2019-07-21	end_session
4	3	2019-06-25	open_session
4	3	2019-06-25	end_session

Output:

day	active_users
2019-07-20	2
2019-07-21	2

```
SELECT      activity_date AS day, count(distinct user_id) AS active_users
FROM        activity
WHERE       DATEDIFF('2019-07-27', activity_date) <30
GROUP BY    activity_date
```

SQL

Exercise

1142. User Activity for the Past 30 Days II

Table: Activity

Column Name	Type
user_id	int
session_id	int
activity_date	date
activity_type	enum

There is no primary key for this table, it may have duplicate rows.

The activity_type column is an ENUM of type ('open_session', 'end_session', 'scroll_down', 'send_message').

The table shows the user activities for a social media website.

Note that each session belongs to exactly one user.

Write an SQL query to find the average number of sessions per user for a period of 30 days ending 2019-07-27 inclusively, rounded to 2 decimal places. The sessions we want to count for a user are those with at least one activity in that time period.

Input:

Activity table:

user_id	session_id	activity_date	activity_type
1	1	2019-07-20	open_session
1	1	2019-07-20	scroll_down
1	1	2019-07-20	end_session
2	4	2019-07-20	open_session
2	4	2019-07-21	send_message
2	4	2019-07-21	end_session
3	2	2019-07-21	open_session
3	2	2019-07-21	send_message
3	2	2019-07-21	end_session
4	3	2019-06-25	open_session
4	3	2019-06-25	end_session

Output:

average_sessions_per_user
1.33

SQL

Answer

Input:

Activity table:

user_id	session_id	activity_date	activity_type
1	1	2019-07-20	open_session
1	1	2019-07-20	scroll_down
1	1	2019-07-20	end_session
2	4	2019-07-20	open_session
2	4	2019-07-21	send_message
2	4	2019-07-21	end_session
3	2	2019-07-21	open_session
3	2	2019-07-21	send_message
3	2	2019-07-21	end_session
4	3	2019-06-25	open_session
4	3	2019-06-25	end_session

Output:

average_sessions_per_user
1.33

```
SELECT      IFNULL(ROUND(COUNT(DISTINCT session_id)/COUNT(DISTINCT user_id), 2),0.00)
            AS average_sessions_per_user
FROM        Activity
WHERE       activity_date BETWEEN '2019-06-28' AND '2019-07-27'
```

SQL

Exercise

1173. Immediate Food Delivery I

Table: Delivery

Column Name	Type
delivery_id	int
customer_id	int
order_date	date
customer_pref_delivery_date	date

delivery_id is the primary key of this table. The table holds information about food delivery to customers that make orders at some date and specify a preferred delivery date (on the same order date or after it).

If the customer's preferred delivery date is the same as the order date, then the order is called immediate; otherwise, it is called scheduled.

Write an SQL query to find the percentage of immediate orders in the table, rounded to 2 decimal places.

Input:

Delivery table:

delivery_id	customer_id	order_date	customer_pref_delivery_date
1	1	2019-08-01	2019-08-02
2	5	2019-08-02	2019-08-02
3	1	2019-08-11	2019-08-11
4	3	2019-08-24	2019-08-26
5	4	2019-08-21	2019-08-22
6	2	2019-08-11	2019-08-13

Output:

immediate_percentage
33.33

SQL

Answer

Input:

Delivery table:

delivery_id	customer_id	order_date	customer_pref_delivery_date
1	1	2019-08-01	2019-08-02
2	5	2019-08-02	2019-08-02
3	1	2019-08-11	2019-08-11
4	3	2019-08-24	2019-08-26
5	4	2019-08-21	2019-08-22
6	2	2019-08-11	2019-08-13

Output:

immediate_percentage
33.33

```
SELECT      ROUND(100*AVG(order_date = customer_pref_delivery_date), 2)
            AS immediate_percentage
FROM        Delivery;
```

SQL

Exercise

176. Second Highest Salary

Table: Employee

```
+-----+-----+
| Column Name | Type |
+-----+-----+
| id          | int  |
| salary      | int  |
+-----+-----+
```

id is the primary key column for this table.

Each row of this table contains information about the salary of an employee.

Write an SQL query to report the second highest salary from the Employee table. If there is no second highest salary, the query should report null.

Input:

Employee table:

```
+-----+-----+
| id | salary |
+-----+-----+
| 1  | 100    |
| 2  | 200    |
| 3  | 300    |
+-----+-----+
```

Output:

```
+-----+
| SecondHighestSalary |
+-----+
| 200                  |
+-----+
```

SQL

Answer

Input:

Employee table:

id	salary
1	100
2	200
3	300

Output:

SecondHighestSalary
200

SELECT

```
(SELECT      DISTINCT Salary
FROM        Employee
ORDER BY    Salary DESC
LIMIT 1 OFFSET 1) AS SecondHighestSalary
```


SQL

Exercise

612. Shortest Distance in a Plane

Table: Point2D

Column Name Type	
x	int
y	int

(x, y) is the primary key column for this table.

Each row of this table indicates the position of a point on the X-Y plane.

The distance between two points $p1(x1, y1)$ and $p2(x2, y2)$ is $\sqrt{(x2 - x1)^2 + (y2 - y1)^2}$.

Write an SQL query to report the shortest distance between any two points from the Point2D table. Round the distance to two decimal points.

Input:
Point2D table:

x y	
-1	-1
0	0
-1	-2

Output:

shortest
1.00

SQL

Answer

Input:
Point2D table:

x	y
-1	-1
0	0
-1	-2

Output:

shortest
1.00

```
SELECT      p1.x, p1.y, p2.x, p2.y,
            Sqrt((Pow(p1.x - p2.x, 2) + Pow(p1.y - p2.y, 2))) AS distance
FROM        point_2d p1
JOIN        point_2d p2
ON          p1.x != p2.x OR p1.y != p2.y
```

```
SELECT      Round(Sqrt(Min((Pow(p1.x - p2.x, 2) + Pow(p1.y - p2.y, 2)))), 2)
            AS shortest
FROM        point_2d p1
JOIN        point_2d p2
ON          p1.x != p2.x OR p1.y != p2.y
```

SQL

Exercise

1107. New Users Daily Count

Table: Traffic

Column Name	Type
user_id	int
activity	enum
activity_date	date

There is no primary key for this table, it may have duplicate rows.

The activity column is an ENUM type of ('login', 'logout', 'jobs', 'groups', 'homepage').

Write an SQL query to reports for every date within at most 90 days from today, the number of users that logged in for the first time on that date. Assume today is 2019-06-30.

Return the result table in any order.

Input:

Traffic table:

user_id	activity	activity_date
1	login	2019-05-01
1	homepage	2019-05-01
1	logout	2019-05-01
2	login	2019-06-21
2	logout	2019-06-21
3	login	2019-01-01
3	jobs	2019-01-01
3	logout	2019-01-01
4	login	2019-06-21
4	groups	2019-06-21
4	logout	2019-06-21
5	login	2019-03-01
5	logout	2019-03-01
5	login	2019-06-21
5	logout	2019-06-21

Output:

login_date	user_count
2019-05-01	1
2019-06-21	2

SQL

Answer

```
SELECT login_date, count(user_id) AS user_count
FROM (SELECT user_id, min(activity_date) AS login_date
      FROM Traffic
      WHERE activity = 'login'
      GROUP BY user_id) t
WHERE datediff('2019-06-30', login_date) <= 90
GROUP BY login_date
```

```
SELECT login_date, count(user_id) AS user_count
FROM (SELECT user_id, min(activity_date) AS login_date
      FROM Traffic
      WHERE activity = 'login'
      GROUP BY user_id) t
WHERE login_date between date_add('2019-06-30',
interval -90 day) and '2019-06-30'
GROUP BY login_date
```

Input:

Traffic table:

user_id	activity	activity_date
1	login	2019-05-01
1	homepage	2019-05-01
1	logout	2019-05-01
2	login	2019-06-21
2	logout	2019-06-21
3	login	2019-01-01
3	jobs	2019-01-01
3	logout	2019-01-01
4	login	2019-06-21
4	groups	2019-06-21
4	logout	2019-06-21
5	login	2019-03-01
5	logout	2019-03-01
5	login	2019-06-21
5	logout	2019-06-21

Output:

login_date	user_count
2019-05-01	1
2019-06-21	2

SQL

Exercise

1112. Highest Grade For Each Student

Table: Enrollments

Column Name	Type
student_id	int
course_id	int
grade	int

(student_id, course_id) is the primary key of this table.

Write a SQL query to find the highest grade with its corresponding course for each student. In case of a tie, you should find the course with the smallest course_id.

Return the result table ordered by student_id in ascending order.

Input:

Enrollments table:

student_id	course_id	grade
2	2	95
2	3	95
1	1	90
1	2	99
3	1	80
3	2	75
3	3	82

Output:

student_id	course_id	grade
1	2	99
2	2	95
3	3	82

SQL

Answer

Input:

Enrollments table:

student_id	course_id	grade
2	2	95
2	3	95
1	1	90
1	2	99
3	1	80
3	2	75
3	3	82

Output:

student_id	course_id	grade
1	2	99
2	2	95
3	3	82

```
SELECT      student_id, MIN(course_id) AS "course_id", grade AS "grade"
FROM        Enrollments
WHERE       (student_id, grade) IN
            (SELECT      student_id, max(grade)
              FROM        Enrollments
              GROUP BY    student_id)
GROUP BY    student_id
ORDER BY    student_id
```